

# Using MobileNet

## Loading the MobileNet Model

Freeze all layers except the top 4, as we'll only be training the top 4

```
In [1]: from keras.applications import MobileNet

# MobileNet was designed to work on 224 x 224 pixel input images sizes
img_rows, img_cols = 224, 224

# Re-loads the MobileNet model without the top or FC layers
MobileNet = MobileNet(weights = 'imagenet',
                      include_top = False,
                      input_shape = (img_rows, img_cols, 3))

# Here we freeze the last 4 layers
# Layers are set to trainable as True by default
for layer in MobileNet.layers:
    layer.trainable = False

# Let's print our layers
for (i, layer) in enumerate(MobileNet.layers):
    print(str(i) + " " + layer.__class__.__name__, layer.trainable)
```

Using TensorFlow backend.

```
0 InputLayer False
1 ZeroPadding2D False
2 Conv2D False
3 BatchNormalization False
4 ReLU False
5 DepthwiseConv2D False
6 BatchNormalization False
```

3 BatchNormalization False  
4 ReLU False  
5 DepthwiseConv2D False  
6 BatchNormalization False  
7 ReLU False  
8 Conv2D False  
9 BatchNormalization False  
10 ReLU False  
11 ZeroPadding2D False  
12 DepthwiseConv2D False  
13 BatchNormalization False  
14 ReLU False  
15 Conv2D False  
16 BatchNormalization False  
17 ReLU False  
18 DepthwiseConv2D False  
19 BatchNormalization False  
20 ReLU False  
21 Conv2D False  
22 BatchNormalization False  
23 ReLU False  
24 ZeroPadding2D False  
25 DepthwiseConv2D False  
26 BatchNormalization False  
27 ReLU False  
28 Conv2D False  
29 BatchNormalization False  
30 ReLU False  
31 DepthwiseConv2D False  
32 BatchNormalization False  
33 ReLU False  
34 Conv2D False  
35 BatchNormalization False  
36 ReLU False  
37 ZeroPadding2D False  
38 DepthwiseConv2D False  
39 BatchNormalization False  
40 ReLU False

43 ReLU False  
44 DepthwiseConv2D False  
45 BatchNormalization False  
46 ReLU False  
47 Conv2D False  
48 BatchNormalization False  
49 ReLU False  
50 DepthwiseConv2D False  
51 BatchNormalization False  
52 ReLU False  
53 Conv2D False  
54 BatchNormalization False  
55 ReLU False  
56 DepthwiseConv2D False  
57 BatchNormalization False  
58 ReLU False  
59 Conv2D False  
60 BatchNormalization False  
61 ReLU False  
62 DepthwiseConv2D False  
63 BatchNormalization False  
64 ReLU False  
65 Conv2D False  
66 BatchNormalization False  
67 ReLU False  
68 DepthwiseConv2D False  
69 BatchNormalization False  
70 ReLU False  
71 Conv2D False  
72 BatchNormalization False  
73 ReLU False  
74 ZeroPadding2D False  
75 DepthwiseConv2D False  
76 BatchNormalization False  
77 ReLU False  
78 Conv2D False  
79 BatchNormalization False  
80 ReLU False  
81 DepthwiseConv2D False  
82 BatchNormalization False

## Let's make a function that returns our FC Head

```
In [2]: def lw(bottom_model, num_classes):  
        """creates the top or head of the model that will be  
        placed ontop of the bottom layers"""  
  
        top_model = bottom_model.output  
        top_model = GlobalAveragePooling2D()(top_model)  
        top_model = Dense(1024,activation='relu')(top_model)  
        top_model = Dense(1024,activation='relu')(top_model)  
        top_model = Dense(512,activation='relu')(top_model)  
        top_model = Dense(num_classes,activation='softmax')(top_model)  
        return top_model
```

## Let's add our FC Head back onto MobileNet

```
In [3]: from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, GlobalAveragePooling2D
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D
from keras.layers.normalization import BatchNormalization
from keras.models import Model

# Set our class number to 3 (Young, Middle, Old)
num_classes = 2

FC_Head = lw(MobileNet, num_classes)

model = Model(inputs = MobileNet.input, outputs = FC_Head)

print(model.summary())
```

Model: "model\_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 224, 224, 3)	0
conv1_pad (ZeroPadding2D)	(None, 225, 225, 3)	0
conv1 (Conv2D)	(None, 112, 112, 32)	864
conv1_bn (BatchNormalization)	(None, 112, 112, 32)	128
conv1_relu (ReLU)	(None, 112, 112, 32)	0
conv_dw_1 (DepthwiseConv2D)	(None, 112, 112, 32)	288
conv_dw_1_bn (BatchNormaliza	(None, 112, 112, 32)	128
conv_dw_1_relu (ReLU)	(None, 112, 112, 32)	0



## Loading our Own Dataset

```
In [4]: from keras.preprocessing.image import ImageDataGenerator

train_data_dir = 'mydata/train/'
validation_data_dir = 'mydata/validation/'

# Let's use some data augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=45,
    width_shift_range=0.3,
    height_shift_range=0.3,
    horizontal_flip=True,
    fill_mode='nearest')

validation_datagen = ImageDataGenerator(rescale=1./255)

# set our batch size (typically on most mid tier systems we'll use 16-32)
batch_size = 32

train_generator = train_datagen.flow_from_directory(
    train_data_dir,
    target_size=(img_rows, img_cols),
    batch_size=batch_size,
    class_mode='categorical')

validation_generator = validation_datagen.flow_from_directory(
    validation_data_dir,
    target_size=(img_rows, img_cols),
    batch_size=batch_size,
    class_mode='categorical')
```

Found 108 images belonging to 2 classes.

Found 26 images belonging to 2 classes.

## Training out Model

- Note we're using checkpointing and early stopping

```
In [5]: from keras.optimizers import RMSprop
        from keras.callbacks import ModelCheckpoint, EarlyStopping

        checkpoint = ModelCheckpoint("human_dataset.h5",
                                    monitor="val_loss",
                                    mode="min",

                                    save_best_only = True,
                                    verbose=1)

        earllystop = EarlyStopping(monitor = 'val_loss',
                                   min_delta = 0,
                                   patience = 3,
                                   verbose = 1,
                                   restore_best_weights = True)

        # we put our call backs into a callback list
        callbacks = [earllystop, checkpoint]

        # We use a very small learning rate
        model.compile(loss = 'categorical_crossentropy',
                      optimizer = RMSprop(lr = 0.001),
                      metrics = ['accuracy'])

        # Enter the number of training and validation samples here
        nb_train_samples = 200
        nb_validation_samples = 50

        # We only train 5 EPOCHS
        epochs = 5
        batch_size = 16
```

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch = nb_train_samples // batch_size,  
    epochs = epochs,  
    callbacks = callbacks,  
    validation_data = validation_generator,  
    validation_steps = nb_validation_samples // batch_size)
```

Epoch 1/5

12/12 [=====] - 42s 3s/step - loss: 6.4891 - accuracy: 0.5123 - val\_loss: 0.7477 - val\_accuracy: 0.3846

Epoch 00001: val\_loss improved from inf to 0.74767, saving model to human\_dataset.h5

Epoch 2/5

12/12 [=====] - 38s 3s/step - loss: 0.6976 - accuracy: 0.5617 - val\_loss: 0.6807 - val\_accuracy: 0.6538

Epoch 00002: val\_loss improved from 0.74767 to 0.68067, saving model to human\_dataset.h5

Epoch 3/5

12/12 [=====] - 38s 3s/step - loss: 0.7015 - accuracy: 0.5833 - val\_loss: 0.6807 - val\_accuracy: 0.3462

Epoch 00003: val\_loss did not improve from 0.68067

Epoch 4/5

12/12 [=====] - 38s 3s/step - loss: 0.8123 - accuracy: 0.6265 - val\_loss: 0.8293 - val\_accuracy: 0.5385

Epoch 00004: val\_loss did not improve from 0.68067

Epoch 5/5

12/12 [=====] - 37s 3s/step - loss: 0.6606 - accuracy: 0.7068 - val\_loss: 1.1875 - val\_accuracy: 0.6538

Restoring model weights from the end of the best epoch

Epoch 00005: val\_loss did not improve from 0.68067

Epoch 00005: early stopping



## Loading our classifier

```
In [6]: from keras.models import load_model

classifier = load_model('human_dataset.h5')
```

## Testing our classifier on some test images

```
In [7]: import os
import cv2
import numpy as np
from os import listdir
from os.path import isfile, join

human_dict = { "[0]": "Akshay Kumar",
               "[1]": "Salman Khan" }

human_dict_n = { "akshaykumar": "Akshay Kumar",
                 "salmankhan": "Salman Khan" }

def draw_test(name, pred, im):
    person = human_dict[str(pred)]
    BLACK = [0,0,0]
    expanded_image = cv2.copyMakeBorder(im, 80, 0, 0, 100 ,cv2.BORDER_CONSTANT,value=BLACK)
    cv2.putText(expanded_image, person, (20, 60) , cv2.FONT_HERSHEY_SIMPLEX,1, (0,0,255), 2)
    cv2.imshow(name, expanded_image)

def getRandomImage(path):
    """function loads a random images from a random folder in our test path """
    folders = list(filter(lambda x: os.path.isdir(os.path.join(path, x)), os.listdir(path)))
    random_directory = np.random.randint(0,len(folders))
    path_class = folders[random_directory]
    print("Class - " + human_dict_n[str(path_class)])
    file_path = path + path_class
```

```

path_class = folders[random_directory]
print("Class - " + human_dict_n[str(path_class)])
file_path = path + path_class
file_names = [f for f in listdir(file_path) if isfile(join(file_path, f))]
random_file_index = np.random.randint(0, len(file_names))
image_name = file_names[random_file_index]
return cv2.imread(file_path+"/"+image_name)

for i in range(0,10):
    input_im = getRandomImage("mydata/validation/")
    input_original = input_im.copy()
    input_original = cv2.resize(input_original, None, fx=0.5, fy=0.5, interpolation = cv2.INTER_LINEAR)

    input_im = cv2.resize(input_im, (224, 224), interpolation = cv2.INTER_LINEAR)
    input_im = input_im / 255.
    input_im = input_im.reshape(1,224,224,3)

    # Get Prediction
    res = np.argmax(classifier.predict(input_im, 1, verbose = 0), axis=1)

    # Show image with predicted class
    draw_test("Prediction", res, input_original)
    cv2.waitKey(1000)

cv2.destroyAllWindows()

```

```

Class - Salman Khan
Class - Akshay Kumar
Class - Akshay Kumar
Class - Salman Khan
Class - Salman Khan
Class - Akshay Kumar
Class - Akshay Kumar
Class - Salman Khan
Class - Akshay Kumar
Class - Akshay Kumar

```